

---

# Scam Detection + Prank Virus

By: Lanz Dungo, Matthew Wolz, Umair Iqbal

---

# Spam Email Classification Using Machine Learning

## Goal:

- Predict whether an email is *spam* or *not spam* using text-based features

## Dataset:

- Two columns:
  1. Label (0 = not spam, 1 = spam)
  2. Text

## Approach:

1. Feature Engineer text-based signals
2. Feed into a Random Forest
3. Feature Explanations

**Question:** *What features contribute most to whether an email is spam or not?*



# Engineered Features & Model Used

## Extracted:

1. Keyword indicators
  - "Free" → has\_free
  - "Cash" → has\_cash
  - "Link" → has\_link
  - "Urgent" → has\_urgent
2. Structure variables
  - Word count → num\_words
  - Question marks → num\_questions
3. Sentiment polarity

Trained a Random Forest Classifier (92% accuracy)

- Optuna hyperparameter tuning

Spam Label	Number of Words	Contains "Free"
0	20	0
0	6	0
1	28	1
0	11	0
0	13	0
1	32	0
0	16	0
0	26	0
1	26	0
1	29	1

Note: the full dataset has 21 columns

# What Predicts Spam?

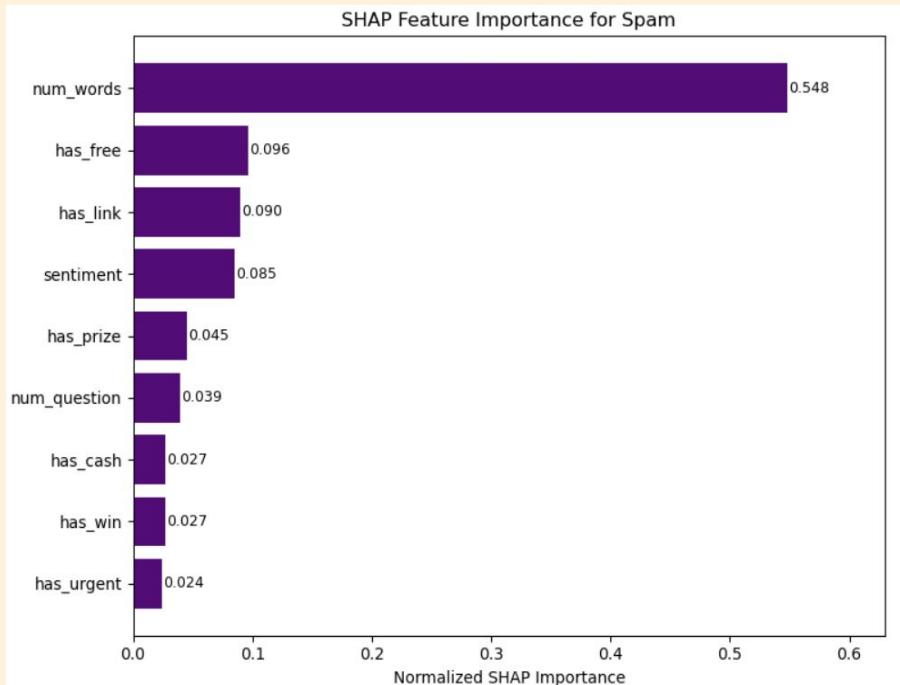
## SHAP values:

- Feature ranking
- Importance values

What variables does the model deem as important?

- num\_words (~55%)
- has\_free
- has\_link
- sentiment

SHAP values → what **signals** do the model consider “spam-like”?



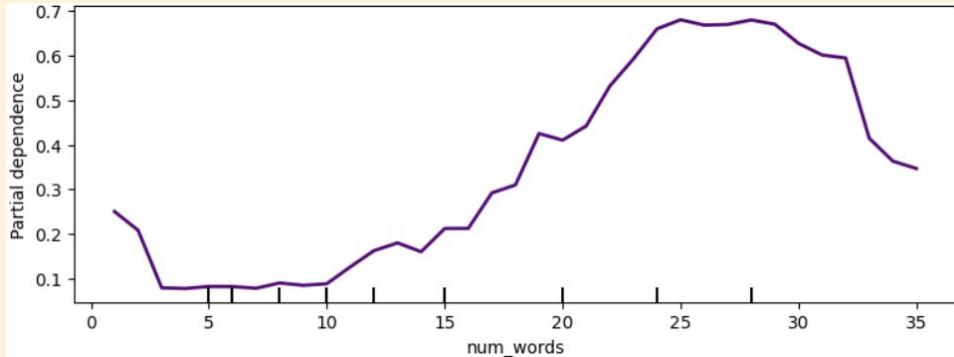
# How Features Affect Spam Probability

## Partial Dependence Plots (PDPs)

- Show how changing a single feature impacts spam probability

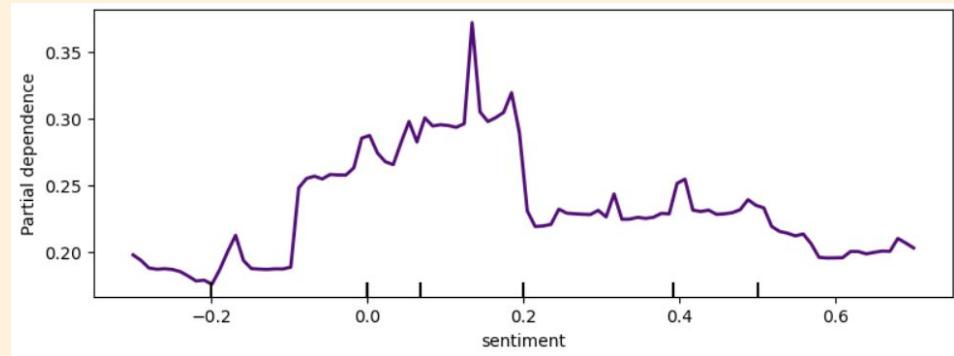
### `num_words`

- Spam probability rises sharply after ~15-20 words in length



### Sentiment

- Neutral/slightly positive sentiment increases spam likelihood



# Key Takeaways

## Spam emails often

- Are longer
- Contain promotional keywords (“free”, “link”, “prize”, “winner”)
- Use upbeat/positive sentiment

The **random forest** model achieves strong performance

- 92% overall accuracy
- Explainable by SHAP values + PDPs

Leads to better knowledge of what a spam email **contains**

